

202A Final Project Code

Stella Huang

12/14/2020

Load Libraries

```
library(tidyverse)
library(dplyr)
library(lubridate)
library(bigmemory)
library(biganalytics)
library(bigtabulate)
library(caret)
library(corrplot)
library(anytime)
library(klaR)
library(devtools)
library(xgboost)
library(gbm)
library(leaps)
```

1. Data Cleaning

```
# load data
data <- read.csv("full_listings.csv") # there are 30533 listings

#####
### data cleaning ###
#####

# remove unnecessary variables
colnames(data)
col_drops <- c("last_scraped", "name", "description", "listing_url", "scrape_id", "neighborhood_overview",
               "host_name", "host_url", "picture_url", "host_about", "host_thumbnail_url", "host_picture_urls",
               "host_location", "host_neighbourhood", "host_has_profile_pic", "neighbourhood_group_clean",
               "calendar_last_scraped", "license", "has_availability", "host_response_time", "neighbourhood_cleanscore")

data <- data[ , !(names(data) %in% col_drops)]

# remove listings wit no rating
data <- data[!is.na(data$review_scores_rating),]
```

data

```
#####  
### text conversion / feature engineering ###  
#####  
  
### PRICE PROCESSING  
# change price to numeric  
data$price <- as.numeric(substr(data$price,2,nchar(data$price)-3))  
  
### HOST INFO  
# change percentages to numeric  
data$host_response_rate[data$host_response_rate=="N/A"] <- "0%"  
data$host_acceptance_rate[data$host_acceptance_rate=="N/A"] <- "0%"  
data$host_response_rate <- as.numeric(substr(data$host_response_rate,1,nchar(data$host_response_rate)-1))  
data$host_acceptance_rate <- as.numeric(substr(data$host_acceptance_rate,1,nchar(data$host_acceptance_rate)-1))  
# count number of host verification documents and amenities  
data <- data %>% mutate(num_host_verif = str_count(host_verifications, "\\")/2, num_amenities=str_count(host_verifications, "\\amenities"))  
  
### BATHROOM COUNT CONVERSION  
# add bathroom num  
# get first element, which is the number of bathrooms  
data$bathrooms <- unlist(lapply(str_split(data$bathrooms_text, " "), ~[[, 1]))  
# check which entries are half baths, and replace with 0  
data[data$bathrooms_text %in% c("Half-bath", "Private half-bath", "Shared half-bath"),]$bathrooms <- 0  
# create variable to indicate if private or not (0:no, 1:yes)  
data$is_private_bath <- ifelse(str_detect(data$bathrooms_text, "private"), 1, 0)  
# convert to numeric  
data$bathrooms <- as.numeric(data$bathrooms)  
  
### ROOM TYPE CONVERSION  
# check room type  
table(data$room_type)  
# assign number encodings to room type  
room_type2 <- data$room_type  
room_type2[room_type2 == "Shared room"] <- 1  
room_type2[room_type2 == "Private room"] <- 2  
room_type2[room_type2 == "Hotel room"] <- 3  
room_type2[room_type2 == "Entire home/apt"] <- 4  
data$room_type2 <- as.numeric(room_type2)  
  
### MISSING VALUES  
data$bedrooms[is.na(data$bedrooms)] <- 0  
data$beds[is.na(data$beds)] <- 0  
data$reviews_per_month[is.na(data$reviews_per_month)] <- 0
```

```

### TRUE/FALSE CONVERSION
# find such col
binary_feat <- c("host_is_superhost", "host_identity_verified", "instant_bookable")
# replace with 0/1 and convert to numeric
binary_feat_data <- data[names(data) %in% binary_feat] %>% mutate_all(funs(str_replace(., "t", "1")))
# replace with 0/1 and convert to numeric
) %>% mutate_all(funs(str_replace(., "f", "0"))) %>%
data[names(data) %in% binary_feat] <- binary_feat_data

### DATE CONVERSION
data %>% select_if(is.character)
date_feat <- c("host_since", "first_review", "last_review")
# convert to date format
for(i in date_feat){
  data[,i] <- as.Date.character(mdy(data[,i]))
}

# calculate durations
data$host_years <- interval(data$host_since, lubridate::ymd("2020-11-23")) / years(1)
data$last_review_date <- interval(data$last_review, lubridate::ymd("2020-11-23")) / years(1)

### MAP NEIGHBORHOOD TO REGION
regions <- read.csv("la_regions.csv")
table(regions$NAME)
# find unmapped neighborhoods and replace names
data[!data$neighbourhood_cleaned %in% regions$NAME,]
data$neighbourhood_cleaned <- str_replace(data$neighbourhood_cleaned, "La Canada Flintridge", "La Cañada Flintridge")
data$neighbourhood_cleaned <- str_replace(data$neighbourhood_cleaned, "East Whittier", "Whittier")
sum(!data$neighbourhood_cleaned %in% regions$NAME)==0
# add regions
data$region <- regions[match(data$neighbourhood_cleaned, regions$NAME),]$REGION
# combine angeles forest w/ antelope valley b/c there's only 9 obs in angeles forest
data$region <- str_replace(data$region, "Angeles Forest", "Antelope Valley")
table(data$region)

```

2. Remove Correlated Features

```

#####
### remove unneeded features ###
#####
var_to_remove <- c("host_verifications", "amenities", "first_review", "last_review", "host_since", "bathrooms")
data <- data[!names(data) %in% var_to_remove]

#####
### feature correlation ###
#####
feat_leaveout <- c("id", "host_id", "latitude", "longitude")
# find correlated feats
cor_matrix <- cor((data[!names(data) %in% feat_leaveout] %>% select_if(is.numeric)), use="complete.obs")

```

```

cor_matrix[upper.tri(cor_matrix)] <- 0 #no symmetry => redundant pairs
diag(cor_matrix) <- 0
# create dataframe of cor var
cor_df <- as.data.frame(as.table(cor_matrix))
# filter with 0.75+ correlation
correlated_var <- cor_df %>% filter(Freq >= 0.75) %>% filter(Var1!=Var2)%>% mutate_if(is.factor, as.character)
# feat_to_keep <- cor_df %>% filter(Freq >= 0.75) %>% filter(Var1!=Var2) %>% select(Var2) %>% pull()

# selected some feat to keep
correlated_var
feat_to_keep <- c("review_scores_rating", "host_listings_count", "bedrooms", "beds",
                 "minimum_nights", "maximum_nights_avg_ntm", "availability_60",
                 "review_scores_rating", "reviews_per_month")
corfeat_to_remove <- unique(c(correlated_var$Var1, correlated_var$Var2)[!c(correlated_var$Var1, correlated_var$Var2) %in% feat_to_keep])
# remove correlated var
corfeat_to_remove
data <- data[!names(data) %in% corfeat_to_remove]

```

2.5 Count listings nearby

```

# use complete data only
data <- subset(data, select = -c(neighbourhood_cleansed, property_type, room_type)) # remove more var
data <- data[complete.cases(data),]

## DISTANCE BASED METRIC
# wanna find the number of listings within 2 mile radius
library(geosphere)
dt.haversine <- function(lat_from, lon_from, lat_to, lon_to, r = 6378137){
  radians <- pi/180
  lat_to <- lat_to * radians
  lat_from <- lat_from * radians
  lon_to <- lon_to * radians
  lon_from <- lon_from * radians
  dLat <- (lat_to - lat_from)
  dLon <- (lon_to - lon_from)
  a <- (sin(dLat/2)^2) + (cos(lat_from) * cos(lat_to)) * (sin(dLon/2)^2)
  return(2 * atan2(sqrt(a), sqrt(1 - a)) * r)
}

num_proximity <- c()

for(i in c(1:nrow(data))){
  # for(i in c(1:10)){
  samp1 <- as.matrix(data[i,c("longitude", "latitude")])
  other_samp <- as.matrix(data[-i,c("longitude", "latitude")])
  pair_distances <- distHaversine(samp1, other_samp)/1609
  # keep samples within 3 miles
  close_listings <- sum(pair_distances<=3)
  num_proximity[i] <- close_listings
}

```

```

summary(num_proximity)
hist(num_proximity, breaks=100)
data$num_proximity <- num_proximity

# ## 2D KERNEL REGRESSION
library(splancs)
library(maps)
latlon_bw <- sqrt(bw.nrd0(data$longitude)^2+bw.nrd0(data$latitude)^2)
latlon <- as.points(data$longitude, data$latitude)

boundary <- matrix(c(range(data$longitude)[1]-0.1,range(data$latitude)[1]-0.1,
                      range(data$longitude)[2]+0.1,range(data$latitude)[1]-0.1,
                      range(data$longitude)[2]+0.1,range(data$latitude)[2]+0.1,
                      range(data$longitude)[1]-0.1,range(data$latitude)[2]+0.1,
                      range(data$longitude)[1]-0.1,range(data$latitude)[1]-0.1),ncol=2,byrow=T)

latlon_dens <- kernel2d(latlon, boundary, latlon_bw)
par(mfrow=c(1,2))
image(latlon_dens, col=gray((64:20)/64),xlab="Latitude",ylab="Longitude", main="2D Kernel Smoothing of 1
points(latlon, col='red')
map('county', 'California', add=T)

```

3. Exploratory Analysis

```

data <- read.csv("cleaned_data.csv")
full_data <- read.csv("full_listings.csv", stringsAsFactors = FALSE)
full_data <- full_data[match(data$id,full_data$id),]
data$local_host <- ifelse(str_detect(full_data$host_location, "Los Angeles"), 1, 0)

# convert char to factor
data[sapply(data, is.character)] <- lapply(data[sapply(data, is.character)], as.factor)
data <- model.frame(price~., data=data, drop.unused.levels = TRUE)
# filter region
data <- data %>% filter(region %in% c("Santa Monica Mountains", "Westside", "Central L.A.", "South L.A.

data$num_amenities

ggplot(aes(x=as.factor(data$room_type2), y=log(price), group=as.factor(data$room_type2)), data=data) +
  geom_boxplot(aes(fill = as.factor(data$room_type2))) + facet_grid(.~region) + theme_light() +
  ggtitle("Room Type vs Price") + xlab("room type") +
  scale_fill_discrete(name = "Room Type", labels = c("Shared","Private","Hotel", "Entire apt/home"))

ggplot(aes(x=price), data=data[data$price<600,]) + geom_histogram(bins=100) + facet_grid(.~region) + th

histogram(log(data$price))
histogram(data$price[data$price<700])

```

3.5. Kernel Regression

C Code

```
# include <R.h>
# include <Rmath.h>
void kernel_reg(int *n, double *x, double *y, double *b, int* m,
               double *g2, double *res2)
{
    int i,j;
    double c;
    double est_sum_y;
    double est_sum;

    // loop through each grid index
    for(i=0; i<*m; i++){
        est_sum = 0.0;
        est_sum_y = 0.0;

        // calculate density; loop through each data index
        for(j=0; j<*n; j++){
            c = dnorm(x[j]-g2[i], 0, *b, 0) / *n; //compute kernel at x on grid
            est_sum += c; // denominator
            est_sum_y += y[j] * c; //multiply by y and add; numerator
        }

        res2[i] = est_sum_y / est_sum;
    }
}

# let's examine kernel regression density of price vs longitude
system("R CMD SHLIB kernel.c")
dyn.load("kernel.so")

# declare var
x <- data$longitude
y <- log(data$price)
xgrid <- c(seq(from = range(x)[1], to = range(x)[2], length.out=300))
bw <- bw.nrd(x)

# load function
gaussian_kernel <- function(x, y, bw, xgrid){
    n <- length(x)
    m <- length(xgrid)
    a=.C("kernel_reg", as.integer(n), as.double(x), as.double(y),
        as.double(bw), as.integer(m), as.double(xgrid), y=double(m))
    a$y
}

kernel_est <- gaussian_kernel(x,y,bw,xgrid)
plot(kernel_est~xgrid, type='l')
kernel_est
```

```

# resampling to get estimate
est_matrix <- c()
for(i in c(1:200)){
  new_samples <- sample_n(data, 2000, replace=TRUE)
  new_y <- as.vector(log(new_samples$price))
  new_x <- as.vector(new_samples$longitude)

  new_kernel_est <- gaussian_kernel(new_x, new_y, bw, xgrid)
  est_matrix <- cbind(est_matrix, new_kernel_est)
}

# extract 5th largest and 195th largest: order each row -> get 5th & 195th col of matrix
dim(est_matrix)
ci95 <- t(apply(est_matrix, 1, sort))[, c(10,195)]
colnames(ci95) <-c("lower", "upper")

kernelest_df <- as.data.frame(cbind(xgrid, kernel_est, ci95))

# plot KDE with 95% CI
ggplot(aes(x=xgrid, y=kernel_est), data=kernelest_df) + geom_point() +
  geom_line() + geom_ribbon(data=kernelest_df,aes(ymin=lower,ymax=upper),alpha=0.3) +
  xlab("Longitude") + ylab("Estimated log(Price)") + ggtitle("KDE Price with 95% CI Band") + theme_minimal()

```

4. Train test split

```

# write.csv(data,"cleaned_data.csv")
data <- read.csv("cleaned_data.csv")
full_data <- read.csv("full_listings.csv", stringsAsFactors = FALSE)
full_data <- full_data[match(data$id,full_data$id),]
data$local_host <- ifelse(str_detect(full_data$host_location, "Los Angeles"), 1, 0)

# convert char to factor
data[sapply(data, is.character)] <- lapply(data[sapply(data, is.character)], as.factor)
data <- model.frame(price~., data=data, drop.unused.levels = TRUE)
# filter region
data <- data %>% filter(region %in% c("Santa Monica Mountains", "Westside", "Central L.A.", "South L.A.

# train test split
train_index <- createDataPartition(data$price, p = .7, list = FALSE, times = 1)
train_data <- data[train_index,]
train_y <- train_data$price
train_data <- subset(train_data, select = -c(price, id, host_id))

test_data <- data[-train_index,]
test_y <- test_data$price
test_data <- subset(test_data, select = -c(price, id, host_id))

```

Function to calculate adjusted r^2

```
adj_r2 <- function(r2, n, p){  
  adj = 1-(((1-r2)*(n-1))/(n-p-1))  
  return (adj)  
}
```

5. Linear Model

```
#####  
### BUILDING LINEAR MODEL #####  
#####  
# drop unused levels in train data  
train_data <- model.frame(train_price~., data=train_data, drop.unused.levels = TRUE)  
linearmodel <- lm(train_y~., data=train_data) # original var  
linearmodel <- lm(log(train_y)~., data=train_data) # transformed var  
summary(linearmodel)  
par(mfrow=c(2,2))  
plot(linearmodel)  
mtext(text=expression(bold("Residual Plot for Linear Model, Log(y)")), side = 3, line = -14, outer = TRUE)  
  
# bc <- boxcox(linearmodel)  
# lam1 <- bc$x[which.max(bc$y)]  
# linearmodel <- lm(((train_price^lam1) - 1) / lam1~., data=train_data)  
  
lm_predictions <- predict(linearmodel, test_data)  
postResample(pred = lm_predictions, obs = log(test_y))  
  
# USE CV  
fit.control <- trainControl(method = "repeatedcv", number = 5, repeats = 3)  
set.seed(123)  
# train linear model with CV  
lm_caret <- train(log(train_y)~., train_data, method = "lm", trControl = fit.control)  
summary(lm_caret)  
lm_predictions <- predict(lm_caret, test_data)  
# calculate prediction accuracy  
postResample(pred = lm_predictions, obs = log(test_y))  
adj_r2(postResample(pred = lm_predictions, obs = log(test_y))[2], nrow(test_data), 36)
```

6. Step-wise regression

```
# step wise selection with leaps library  
sw_linearmod <- regsubsets(log(train_y)~., data=train_data, nvmax = 27, method = "seqrep")  
reg_summary <- summary(sw_linearmod)  
par(mfrow = c(2,2))  
plot(reg_summary$rss, xlab = "Number of Variables", ylab = "RSS", type = "l")  
plot(reg_summary$adjr2, xlab = "Number of Variables", ylab = "Adjusted RSq", type = "l")  
  
# We will now plot a red dot to indicate the model with the largest adjusted  $R^2$  statistic.
```



```

# The which.max() function can be used to identify the location of the maximum point of a vector
adj_r2_max = which.max(reg_summary$adjr2) # 11

# The points() command works like the plot() command, except that it puts points
# on a plot that has already been created instead of creating a new plot
points(adj_r2_max, reg_summary$adjr2[adj_r2_max], col = "red", cex = 2, pch = 20)

# We'll do the same for C_p and BIC, this time looking for the models with the SMALLEST statistic
plot(reg_summary$cp, xlab = "Number of Variables", ylab = "Cp", type = "l")
cp_min = which.min(reg_summary$cp) # 10
points(cp_min, reg_summary$cp[cp_min], col = "red", cex = 2, pch = 20)

plot(reg_summary$bic, xlab = "Number of Variables", ylab = "BIC", type = "l")
bic_min = which.min(reg_summary$bic) # 6
points(bic_min, reg_summary$bic[bic_min], col = "red", cex = 2, pch = 20)
mtext(text=expression(bold("Step-Wise Regression Performance")), side = 3, line = -1.5, outer = TRUE)

# choose best number of var to include
# find such coefficients
sw_bestvar <- names(coef(sw_linearmod, 23))

fit.control <- trainControl(method = "repeatedcv", number = 5, repeats = 3)

set.seed(123)
# fit model with best subset selected
lm_caret <- train(log(train_y)~., train_data[,names(train_data) %in% c(sw_bestvar, "region")],
                  method = "lm", trControl = fit.control)
lm_summary <- as.data.frame((exp(coef(summary(lm_caret)))-1)*100)
options(scipen=999)
lm_summary[order(lm_summary$Estimate),]

# calculate accuracy
lm_predictions <- predict(lm_caret, test_data)
postResample(pred = lm_predictions, obs = log(test_y))
adj_r2(postResample(pred = lm_predictions, obs = log(test_y))[2], nrow(test_data), 36)

```

7. Recursive Feature Elimination

```

#####
# ### RFE FOR LINEAR MODEL #####
# #####
ctrl <- rfeControl(functions = lmFuncs,
                   method = "repeatedcv",
                   repeats = 5,
                   number = 3,
                   verbose = FALSE)

lmProfile <- rfe(x=train_data%>%select_if(is.numeric), y=log(train_y), sizes = c(1:20), rfeControl = ctrl)

# rank of variables

```

```

lmProfile$variables
best_rfe_lm <- lmProfile$variables %>% filter(Variables==32)

# sapply(best_rfe_lm, mean)

best_rfe_lm_var <- best_rfe_lm %>% filter(var %in% c("bedrooms", "room_type2", "longitude", "is_private"))

best_rfe_lm_var$var <- str_replace(best_rfe_lm_var$var,
                                   "calculated_host_listings_count_shared_rooms", "host_sharedRooms_cnt")

best_rfe_lm_var$var <- str_replace(best_rfe_lm_var$var, "room_type2", "room type")
ggplot(aes(x=reorder(var, Overall), y=Overall), data=best_rfe_lm_var) + geom_bar(stat="identity", fill="white")
+ coord_flip() + theme_minimal() + geom_text(aes(label=round(Overall,2)), hjust=1.1, size=3.5, col="white")
+ ggtitle("Variable Importance from LM RFE") + xlab("Importance") + ylab("Variable")

```

8. Random Forest

```

#####
### RANDOM FOREST REGRESSION ###
#####
library(randomForest)
library(e1071)

mtry <- round(ncol(train_data)/3)
tuneGrid <- expand.grid(.mtry=mtry)
train.control <- trainControl(method = "repeatedcv", number = 5, repeats=3)

# Train the model
rf_model <- train(log(train_y)~., data=train_data,
                  method = "rf", tuneGrid=tuneGrid,
                  trControl = train.control, importance=T)

# RESULTS
print(rf_model)
#####
# 8597 samples
# 32 predictor
#
# No pre-processing
# Resampling: Cross-Validated (5 fold, repeated 3 times)
# Summary of sample sizes: 6877, 6877, 6877, 6878, 6879, 6878, ...
# Resampling results
#
# RMSE          Rsquared    RMSE SD       Rsquared SD
# 0.3498099    0.7528557    0.007383131  0.01263049
#
# Tuning parameter 'mtry' was held constant at a value of 11
#####

rf_model$finalModel #output below
#####

```

```

# Call:
# randomForest(x = x, y = y, mtry = param$mtry, importance = ..1)
#           Type of random forest: regression
#           Number of trees: 500
# No. of variables tried at each split: 11
#
#           Mean of squared residuals: 0.1196127
#           % Var explained: 75.54
#####

# make prediction
rf_predictions <- predict(rf_model, test_data, type="raw")
postResample(pred = rf_predictions, obs = log(test_y))
adj_r2(postResample(pred = rf_predictions, obs = log(test_y))[2], nrow(train_data), 32)

# VARIABLE IMPORTANCE
write.csv(varImp(rf_model, scale = TRUE)$importance, "price_rfImp.csv")
rf_imp <- read.csv("price_rfImp.csv")

# plot variable importance
ggplot(aes(x=Overall), data=varImp(rf_model, scale = TRUE)) + geom_bar(stat="identity", fill="steelblue") +
  coord_flip() + theme_minimal() + geom_text(aes(label=round(Importance,2)), hjust=1.1, size=3.5, col="white") +
  ggtitle("Variable Importance from Random Forest") + xlab("Importance") + ylab("Variable")

ggplot(aes(x=reorder(X, Overall), y=Overall), data=(rf_imp %>% arrange(desc(Overall)))[1:7,]) +
  geom_bar(stat="identity", fill="steelblue") + coord_flip() + theme_minimal() +
  geom_text(aes(label=round(Overall,2)), hjust=1.1, size=3.5, col="white") + ggtitle("Variable Importance from Random Forest") +
  xlab("Importance") + ylab("Variable")

```